

# Data Management and Databases

Spring 2021

Jean-Claude Graf

July 16, 2021

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Terminology . . . . .	3
1.2	Different Models . . . . .	3
<b>2</b>	<b>Relational Model</b>	<b>4</b>
2.1	Schema . . . . .	4
2.1.1	Instance . . . . .	4
2.2	Key . . . . .	4
<b>3</b>	<b>Algebras</b>	<b>5</b>
<b>4</b>	<b>Relational Algebra</b>	<b>5</b>
<b>5</b>	<b>Relational Calculus</b>	<b>6</b>
<b>6</b>	<b>SQL</b>	<b>8</b>
6.1	DDL . . . . .	8
6.2	DML . . . . .	9
6.3	Query Language . . . . .	9
6.4	Tips . . . . .	10
<b>7</b>	<b>Graphical Modelling</b>	<b>11</b>
7.1	Conceptual Modelling (ER-Diagram) . . . . .	11
7.2	Logical Modelling . . . . .	13
<b>8</b>	<b>Integrity Constraints</b>	<b>14</b>
8.1	Types . . . . .	14
<b>9</b>	<b>Recursive Queries</b>	<b>16</b>
<b>10</b>	<b>NULL</b>	<b>17</b>
<b>11</b>	<b>Views</b>	<b>18</b>
<b>12</b>	<b>Functional Dependency</b>	<b>19</b>
<b>13</b>	<b>Normal Form</b>	<b>22</b>

<b>14 Analytic</b>	<b>25</b>
14.1 Associated Rule Mining . . . . .	25
14.2 Clustering . . . . .	26
14.3 Classification . . . . .	27
<b>15 DMDB System</b>	<b>28</b>
15.1 Storage Hierarchy . . . . .	28
15.2 Disk Manager . . . . .	28
15.2.1 File Layout . . . . .	29
15.2.2 Page Layout . . . . .	30
15.2.3 Tuple layout . . . . .	30
15.3 Buffer Pool Management . . . . .	31
15.4 Access Methods . . . . .	32
15.5 Operator Execution . . . . .	34
15.5.1 Query Optimizer . . . . .	36
15.6 Tuning . . . . .	39
<b>16 Transaction</b>	<b>40</b>
16.1 ACID: . . . . .	40
16.2 Isolation . . . . .	41
16.3 More on Serializable . . . . .	41
16.4 Enforce Isolation . . . . .	42
16.4.1 Locking . . . . .	42
16.4.2 Snapshot Isolation . . . . .	44
<b>17 Recoverability</b>	<b>45</b>
17.1 Definition . . . . .	45
17.2 Write-Ahead Log . . . . .	45
<b>18 Distribution</b>	<b>48</b>

# 1 Introduction

## 1.1 Terminology

- **Database (DB):** collection of data
- **Database Management System (DBMS):** software for maintaining and utilizing a DB
  - ◊ Desired properties:
    - \* **Data Independance:** applications should not know how data are stored
    - \* **Declarative Efficient Data Access:** system should be efficient
    - \* **Transactional Access:** simulate each user that it is the only one interacting with the system
    - \* **Generic Abstraction:** users should not worry about all the issues above
  - ◊ Different flavours exists for different purposes

## 1.2 Different Models

- **Hierarchical Model**
  - ◊ Introduced by IBM
  - ◊ 1968 - today
  - ◊ Hierarchical structure of entities
  - ◊ Imperative query (say how we want it)
    - No data independence
    - No declarative efficient data access
- **Network Model**
  - ◊ 1969
  - ◊ Today in XML and JSON
  - ◊ Network of entities and relations
  - ◊ Imperative query (say how we want it)
    - No data independence
    - No declarative efficient data access
- **Relational Model**
  - ◊ 1969
  - ◊ One of the most popular today
  - ◊ Data stored in tables
  - ◊ Declarative query (say what we want)
  - + Data independence
  - + Declarative efficient data access

## 2 Relational Model

- Knowledge is represented as a *collection of facts*
- Inference is done using *mathematical logic*

### 2.1 Schema

- **Database Schema:**
  - ◊ Set of relation schema
- **Relation Schema:**
  - ◊ “Represented” as a table
  - ◊ Has a name
  - ◊ Contains a set of fields/attributes
  - ◊ Sometimes referred to as *Relation*
  - ◊ Described as  $R(f_1 : D_1, \dots, f_n : D_n)$ 
    - \* **R**: relation name
    - \* **f<sub>i</sub>**: name of field *i*
    - \* **D<sub>i</sub>**: domain of field *i*
- **Field/Attribute:**
  - ◊ “Represented” as a single columns of the table
  - ◊ Has a name
  - ◊ Described by a domain (/type)
- Describes only the header (does not contain any content)
- Is not unique
  - ◊ Different schema have different advantages/disadvantages

#### 2.1.1 Instance

- “Represented” a set of rows in the table
- Set of tuples  $I_R \subseteq D_1 \times \dots \times D_n$  for  $R(f_1 : D_1, \dots, f_n : D_n)$
- **Domain Constraint:** For each field the domain and the schema domain must match
- In practice a DB is a bag and not a set (allows duplicate entries)
  - ◊ But in theory we assume it is a set
- Attributes have no ordering principle, but ordering of attributes and tuples has to match

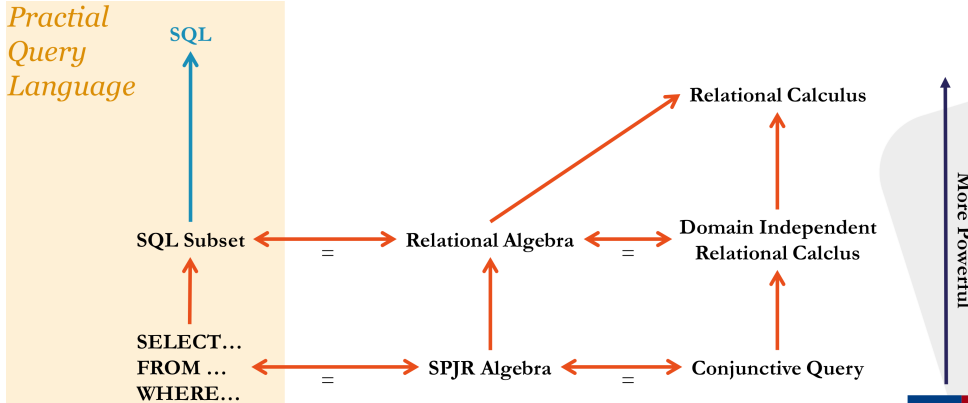
### 2.2 Key

- **Candidate Key:** minimal set of fields that uniquely identify a tuple
- **Primary Key:** one candidate key
  - ◊ Indicated by underlining the field:  $R(\underline{f_1} : D_1, \dots, f_n : D_n)$
  - ◊ Every relation must have one (but in a DB this is not required)
- **Key Constraint:**
  - ◊ The primary key must be unique in each an instance
  - ◊ All valid instances  $I \subseteq D_k \times D_a \times D_b \wedge \forall (k, a, b), (k', a', b') \in I, k = k' \implies (a, b) = (a', b')$

### 3 Algebras

- Overview

*Practical  
Query  
Language*



### 4 Relational Algebra

- Imperative

- ◇ Say how we want it
- ◇ Different ways to implement the same query

- Operators**

- ◇ Create a new relation  $R'$  from one or two given relations  $R_1, R_2$
- ◇ **Union**  $\cup$ :
  - \*  $x \in R_1 \cup R_2 \iff x \in R_1 \vee x \in R_2$
  - \* Schema of both operands must match
- ◇ **Difference**  $-$ :
  - \*  $x \in R_1 - R_2 \iff x \in R_1 \wedge \neg(x \in R_2)$
  - \* Schema of both operands must match
- ◇ **Intersection**  $\cap$ :
  - \*  $x \in R_1 \cap R_2 \iff x \in R_1 \wedge x \in R_2$
  - \* Schema of both operands must match
  - \* Follows from  $R_1 \cap R_2 = R_1 - (R_1 - R_2)$
- ◇ **Selection**  $\sigma$ :
  - \*  $x \in \sigma_c(R) \iff x \in R \wedge c(x) = \text{True}$
  - \* Return tuples which satisfy a given condition  $c$
- ◇ **Projection**  $\pi$ :
  - \*  $\pi_{A_1, \dots, A_n}(R)$  keeps only columns  $A_1, \dots, A_n$
- ◇ **Cartesian Product**  $\times$ :
  - \*  $(x, y) \in R_1 \times R_2 \iff x \in R_1 \wedge y \in R_2$
  - \* If  $R_1$  and  $R_2$  have columns in common renaming is required
  - \* Rarely used in practice
- ◇ **Renaming**  $\rho$ :
  - \*  $\rho_{B_1, \dots, B_n}(R)$  changes the name of the attributes to  $B_1, \dots, B_n$
  - \*  $\rho_S(R)$  changes the names of the attributes of  $R$  to the one of  $S$
- ◇ **Natural Join**  $\bowtie$ :
  - \*  $R_1(A, B) \bowtie R_2(B, C) = \pi_{A, B, C}(\sigma_{R_1.B=R_2.B}(R_1 \times R_2))$ 
    - **No shared attributes:** Equivalent to cross product
    - **All attributes shared:** Equivalent to intersect
  - \* Is associative

- ◇ **Theta Join**  $\bowtie_\theta$ :
  - \*  $R_1 \bowtie_\theta R_2 = \sigma_\theta(R_1 \times R_2)$ 
    - $\theta$  can be any kind of condition
  - \* Flavour
    - **Equi-Join**:  $R_1 \bowtie_{A=B} R_2 = \sigma_{A=B}(R_1 \times R_2)$
- ◇ **Inner Join**:
  - \* Very similar to theta join with the difference that it keeps all matched columns
    - I.e. matching columns are not collapsed into a single one
- ◇ **Outer Join**:
  - \* **Left Outer Join**  $\ltimes$ : Natural join  $\cup$  unmatched tuples from left operand
  - \* **Right Outer Join**  $\rtimes$ : Natural join  $\cup$  unmatched tuples from right operand
  - \* **Full Outer Join**  $\Join$ : Natural join  $\cup$  unmatched tuples from left and right operand
  - \* Unmatched tuples from the left/right are extended with NULL in the columns which do not match
- ◇ **Semi Join**:
  - \* **Left Semi Join**  $\ltimes$ : Tuples from left operand which match with any tuples from right operand
    - $R_1(A_1, \dots, A_n) \ltimes R_2(B_1, \dots, B_2) = \Pi_{A_1, \dots, A_n}(R_1 \bowtie R_2)$
  - \* **Right Semi Join**  $\rtimes$ : Tuples from right operand which match with any tuples from left operand
- ◇ **Relational Division**  $\div$ :
  - \*  $A \div B = C$  where  $C$  is the largest relation such that  $B \times C \subseteq A$
  - \* Inverse of cross product
- **Expression**: Composition of multiple operations
  - No infinite (recursive) queries
- **Bags**
  - ◇ Real DBMS use bag semantic
  - ◇ Can have duplicates
  - ◇ Difficult to extend RA to bags
    - \* Need to add new operators (duplicate elimination)

## 5 Relational Calculus

- Declarative
  - ◇ Say what we want
- More powerful than RA
- **Database Schema**:  $S = (R_1, \dots, R_m)$ ,  $R_i$  is a relation
- **Relation Schema**:  $R(A_1 : D_1, \dots, A_n : D_n)$
- **Domain**:  $\text{dom} = \bigcup_i D_i$ 
  - ◇ Infinite set of constants
- **Instance of Relation**  $R(\mathbf{A}_1 : \mathbf{D}_1, \dots, \mathbf{A}_n : \mathbf{D}_n)$ :  $I_R \subseteq \text{dom}^n$ 
  - ◇  $I_R$  is finite
  - ◇ Set of facts over the relation
- **Instance of DB**  $S(\mathbf{R}_1, \dots, \mathbf{R}_m)$ : Function  $\mathcal{I}$  that maps relation to an instance of that relation
  - ◇  $\mathcal{I} : R_i \mapsto \mathcal{I}(R_i)$
  - ◇ Finite
  - ◇ Set of facts over all relations

- **Query  $Q_\phi$ :** Has the form  $Q_\phi = \{(x_1, \dots, x_k) \mid \phi\}$ 
  - ◊  $\phi$ : First order formula (predicate) with free variables  $x_1, \dots, x_k$
- TODO: Full mathematical definition**
- Output of queries may be infinite
- **Safe:** Query  $Q_\phi$  which is finite for all  $\mathcal{I}$ 
  - ◊ Undecidable problem if query is safe
- **Domain Independent Relational Calculus:** Query whose result does not depend on the interpretation of the relation and not on the domain
- **Active Domain:**  $\text{adom}(Q_\phi, \mathcal{I}) = \text{all constants in } Q_\phi \text{ and } \mathcal{I}$
- **Conjunctive Query:**  $\phi = \exists y_1, \dots, y_l (A_1 \wedge \dots \wedge A_m)$ ,  $Q_\phi = \{(x_1, \dots, x_n) \mid \phi\}$ ,  $A_j$  is an atom
  - ◊ Has many good properties
- **SPJR Algebra:** Relational algebra with only select, project, join and rename operators

## 6 SQL

- Consists of three steps
  - ◊ Define the schema of the tables
  - ◊ Put information into the tables
  - ◊ Query the tables
- SQL is a family of standards
  - ◊ **Data Definition Language (DDL):**
  - ◊ **Data Manipulation Language (DML):**
  - ◊ **Query Language:**
- Released in 1974
- Constantly improved and extended
- Different DB engines implement slightly different standards
  - ◊ Important to choose the right engine for a certain project

### 6.1 DDL

- Defines schema
- Relation schema requires
  - ◊ Name
  - ◊ Set of columns
  - ◊ Type of columns
- Data Types
  - ◊ **char (n):** String of length  $n$ 
    - \* Padded with whitespace to match length
  - ◊ **varchar (n):** String of length  $\leq n$
  - ◊ **integer**
  - ◊ **blob of raw data**
  - ◊ **date**
  - ◊ etc.

- **Create Relation:**

```
CREATE TABLE Professor(  
    PersNr integer,  
    Name varchar (30),  
    Level varchar(2),  
    Room integer,  
    PRIMARY KEY (PersNr));
```

- **Delete Relation:**

```
Drop TABLE Professor;
```

- **Add New Column:**

```
ALTER TABLE Professors ADD COLUMN (age integer);
```

- ◊ Unclear what to insert into preexisting tuples for that relation
- ◊ We can set a default value
  - \* If not provided it keeps the entry empty

- **Drop Column:**

```
ALTER TABLE Professor DROP COLUMN age;
```



## 6.2 DML

- Put and manipulation information
- **Extract, Transform, Load (ETL):** Used to populate DB automatically
  - ◊ Technique for populating DB
  - ◊ **E:** Get data from somewhere
  - ◊ **T:** Bring in right format
  - ◊ **L:** Insert into DB
  - ◊ Manual population is not feasible
- Manual population of DB
  - Error-Prone
  - Slow
    - \* Each query has to be parsed one-by-one
      - Even when all are pretty much equivalent
    - \* Constraints have to be checked for each query

- **Insert:**

```
INSERT INTO Student (PersNr, Name)
VALUES (1111, 'Fred');
```

- **Delete:**

```
DELETE Student
WHERE Semester > 13;
```

- **Update:**

```
UPDATE Student
SET Semester = Semester + 1;
```

- **From CSV:**

```
COPY Professors FROM '/profs.csv' WITH FORAMT csv;
```

## 6.3 Query Language

- **SELECT ... FROM ... WHERE ...**
  - ◊ **FROM:** List of relation whose cross product is taken
  - ◊ **WHERE:** Selection condition
  - ◊ **SELECT:** Projection
- Select uses bag semantic
  - ◊ **SELECT DISTINCT:** Set semantic over all fields
- Every RA expression can be written in *SQL Subset*
- **SQL Subset:**
  - ◊ **Base Query:**  $\rho_{a_1, \dots, a_n}(\Pi_{A_1, \dots, A_n}(\sigma_{P_1 \wedge \dots \wedge P_m}(R_1 \times \dots \times R_k))) \approx$ 

```
SELECT A1 as a1 ... An as an
FROM R1 ... Rk
WHERE P1 AND P2 ... AND Pm;
```
  - ◊ **Union:**  $R_1 \cup R_2$ 

```
(SQL1) UNION (SQL2);
```

    - \* Uses set-like semantic
    - \* **UNION ALL:** used bag semantic

- ◊ **Intersection:**  $R_1 \cap R_2$   
(SQL1) INTERSECT (SQL2);
- ◊ **Difference:**  $R_1 - R_2$   
(SQL1) EXCEPT (SQL2);
- ◊ **Selection:**  $\sigma_c(R)$   
SELECT \* FROM (SQL1) WHERE c;
- ◊ **Projection:**  $\Pi_{A_1, \dots, A_n} R$   
SELECT A1, ..., An FROM (SQL1);
- ◊ **Cross Product:**  $R_1 \times R_2$   
SELECT \* FROM (SQL1), (SQL2);
- ◊ **Rename:**  $\rho_{a,b,c} R$   
SELECT A as a, B as b, C as c FROM (SQL1);
- **Sorting**  
SELECT A, B FROM (SQL1) ORDER BY A DESC, B ASC;
  - ◊ ASC is default
- **Grouping, Aggregation**  
SELECT A, COUNT(\*) FROM (SQL1) GROUP BY A;
  - ◊ Can only project to aggregated fields and aggregation function
  - ◊ **HAVING:** Filter similar to WHERE but for the aggregated field
- **EXISTS, ANY, ALL, SOME:**
  - ◊ Useful things
- **Snapshot Semantics:**
  - ◊ Problematic when deleting (updating) from a relation which we need in a subquery of that relation
  - ◊ Would lead to non-determinism
  - ◊ First all tuples which would get modified are marked
  - ◊ Then the updates are implemented

## 6.4 Tips

- Extract substring by index: `substring(<string>, <base1startindex>, <length>)`
- Operator giving bag:
  - ◊ SELECT column FROM table
  - ◊ SQL1 UNION ALL SQL2
- Operator giving set
  - ◊ SELECT DISTINCT column FROM table
    - \* We can also apply this to multiple columns. In that case the entries are combined and then duplicates filtered.
  - ◊ SQL1 UNION SQL2
- Get current datetime: `NOW()`
- Extract certain field from datetime: `DATE_PART(<desired_field>, <source>)`
  - ◊ `desired_field` can be anything like: `day`, `hour`, `year` etc.
- Data type conversion is done using `<some_field>::<some_type>`
- Casting data type using `CAST(<some_filed> AS <some_type>)`
- Round to  $n$  decimal points: `ROUND(<source>, <n>)`

## 7 Graphical Modelling

- Models an application
- Graphical way to represent entities and their relation
- Consists of three steps
  - ◊ **Conceptual Modeling:** Capture of domains to be represented
    - \* Create diagram from “real world”
    - \* We consider Entity Relation (ER) Model in this course
    - \* Specifies all DB instances that are valid/allowed in our application
  - ◊ **Logical Modeling:** Map concepts to a concrete logical representation
    - \* Convert diagram to table schema
  - ◊ **Physical Modeling:** Implementation in Hardware
    - \* Convert table to bits

### 7.1 Conceptual Modelling (ER-Diagram)

- **Formal Semantics**
  - ◊ Diagram defines valid DB instances
  - ◊ All values we can take  $\mathcal{D} = \mathcal{B} \cup \Delta$ 
    - \*  $\mathcal{B}$ : Concrete values
      - Int, String, Float, etc
    - \*  $\Delta$ : Abstract values
      - Correspond to an entity
  - ◊ **Entity Set  $E$ :** 1-ary Predicate  $E(x)$ 
    - \*  $E(x) = \text{True}$  if  $x$  is of Entity Type  $E$
    - \*  $E^{\mathcal{J}} \subseteq \Delta$
  - ◊ **Attribute  $A$ :** Binary Predicate  $A(x, y)$ 
    - \*  $A(e, a) = \text{True}$  if  $e$  has attribute  $a$
    - \*  $A^{\mathcal{J}} \subseteq \Delta \times \mathcal{B}$
  - ◊  **$n$ -ary Relation  $R$ :**  $n$ -ary Predicate  $R(x_1, \dots, x_n)$ 
    - \*  $R(x_1, \dots, x_n) = \text{True}$  if  $(x_1, \dots, x_n)$  participate in  $R$
    - \*  $R^{\mathcal{J}} \subseteq \Delta^n$
  - ◊ Each subgraph introduces a first-order logic sentence
  - ◊ Entity  $E_1$  and  $E_2$  linked by relation  $R$ 
    - \*  $\forall x_1, x_2 \in \Delta. R(x_1, x_2) \implies E_1(x_1) \wedge E_2(x_2)$
  - ◊ Entity  $E$  with attribute  $A$ 
    - \*  $\forall x, E(x) \implies \underbrace{E^{\mathcal{J}}}_{\text{uniquely exists}} y. A(x, y) \wedge y \in \mathcal{B}$
- **Building Blocks**
  - ◊ **Entity:** Instance of an entity set which is distinguishable from other instances of the same set
  - ◊ **Entity Set:** Set of entities of the same “type”
    - \* Rectangular box
  - ◊ **Attributes:** Properties of a certain entity set
    - \* Round box
  - ◊ **Relationships:** Connection among  $\geq 2$  entity sets
    - \* Rhombus box
    - \* **Roles**
      - Each entity set can have a role in a relation
      - Label lines by the role the entity set is

- ◇ **Key:** Minimal set of attributes which uniquely identify an entity in the entity set
  - \* **Candidate Key:** All possible sets of keys
  - \* **Primary Key:** One selected key
    - Every entity set must have one
    - Underlined
- **Cardinality**
  - ◇ Two main notations
  - ◇ **N/M-Notation**
    - \* **One to One (1/1):**
      - $A$  is in a one to one relationship with  $B$  if:
        - ▷ 1  $A$  entity can only have one relation with a  $B$  entity and
        - ▷ 1  $B$  entity can only have one relation with an  $A$  entity
    - \* **One to Many (1/N):**
      - $A$  is in a one to many relation with  $B$  if:
        - ▷ 1  $A$  can have relationships with multiple  $B$  entities and
        - ▷ 1  $B$  can only have one relation with an  $A$  entity
    - \* **Many to One (N/1):**
    - \* **Many to Many (N/M):**
      - $A$  is in a many to many relation with  $B$  if:
        - ▷ 1  $A$  entity can have relationships with multiple  $B$  entities and
        - ▷ 1  $B$  entity can have relationships with multiple  $A$  entities
  - ◇ **(min, max)-Notation**
    - \* For a relation we give the min and the max value of relations one entity can have
    - \* Stronger than N/M-notation
    - \* \* means infinity
    - \* (min, max) is written in opposite was to N/M
- **Weak Entity**
  - ◇ Some entity relation depends on other entity
    - \* I.e. it is not unique by itself
    - \* Can only be uniquely identified with the main entity
  - ◇ *Weak* entity is the one which depends on another
  - ◇ Indicated by dotted underline
  - ◇ Is a 1/ $M$  relationship
- **Generalisation**
  - ◇ Represent that a entity set is is an instance of another entity set
  - ◇ Entity  $A$  is a entity of  $B$ 
    - \*  $A \subseteq B$
    - \* Draw an arrow from  $A$  to  $B$
    - \*  $A$  shares  $B$ s attributes and primary key
    - \* Are not enforced
      - I.e. possible that  $B \notin A$
    - \* If  $A$  is a $B$  and  $C$  is a $B$  it is possible that  $C \in A$  and  $C \in B$
- There are many other flavours of ER
- **Design Principles**
  - ◇ Model should reflect the application we want to build
  - ◇ Avoid redundancy
  - ◇ Keep it as simple as possible; less entities is better
  - ◇ Entity if the concept has more than one relationship
  - ◇ Attribute if the concept has only one 1:1 relationship

- ◇ Models are large, partition it

## 7.2 Logical Modelling

- Take ER-model and convert to relational model
- Some constraints get lost
- **Steps**
  1. **Entity Sets:** Become relations
  2. **Attributes:** Become attributes of the relation
  3. **Relationship:**
    - ◇ **Without Cardinality Constrain (or N:M):**
      - \* Become relation containing the attributes of all participating relations
      - \* The primary key of the relation are all the primary keys together
    - ◇ **With Cardinality Constraints:**
      - \* Very tricky
      - \* Become relation containing the attributes of all participating relations
      - \* The primary key of the relation are the keys of the entities with which the relation can be uniquely identified. Or the relation gets merged into the table on the *many* side.
    - ◇ **Role:** Can be used to distinguish columns with the same entity type.
      - \* Done by renaming the two columns appropriately
  4. **Weak Entity:**
    - ◇ Can be omitted
    - ◇ The week entity is modeled as a relation on its own with the primary key of the main relation and its own key
  5. **Generalisation:**
    - ◇ Two ways to represent this
    - ◇ Better way depends on application
    - 1) *Child* has its own relation and the *Parent* relation
    - 2) Each *Child* is a full blown relation containing all keys of the *parent* relation
      - Lot of redundant data if entity is multiple child and parent at the same time
      - Cannot constraint that entity is only on of them
- **Rezept**
  1. Convert entries to relations
    - ◇ All attributes of the model get attributes of the relation
    - ◇ All keys of the model get keys of the relation
  2. Convert relations to relations
    - ◇ All attributes of the model get attributes of the relation
    - ◇ All keys of the participating entries are attributes of the relation
    - ◇ Keys are the keys of the entities which uniquely identify the relation
    - ◇ Mark alternative keys
  3. Merge relation if it is  $1 : 1, 1 : N, N : 1$  and it has the same key
  4. Do some other merging
    - ◇ Automatically generates at least 3NF
- Can be done (Semi-) automatically

## 8 Integrity Constraints

- Additional constraint to the key and domain constraint
- Makes sure changes are consistent
- Control the content of the data and its consistency
- Are enforced by the schema
- Can be defined when:
  - ◊ Creating the table (`CREATE table`)
  - ◊ Later (`ALTER table`)
- Checked at `INSERT` as well as `UPDATE`
  - ◊ For foreign key also on `DELETE`
- Check happen at tuple level and not at the semantic of the command
  - ◊ I.e. try to run it and see what happens instead of analysing the query
- Some check may fail or succeed depending on the order of the tuples
  - ◊ We have no influence on this

### 8.1 Types

- **NOT NULL**
  - ◊ Prevents attribute from being `NULL`
  - ◊ **Syntax:** `some_field any_type not null`
- **PRIMARY KEY**
  - ◊ Mark attribute as primary key
  - ◊ Must not be `NULL` and not empty
  - ◊ **Syntax:** `some_field any_type PRIMARY KEY`
  - ◊ If applied to a tuple, all field must not be `NULL`
  - ◊ **Syntax:** `PRIMARY KEY (field1, fiels2)`
- **UNIQUE**
  - ◊ Value must be unique or `NULL`
    - \* In contrast to **PRIMARY KEY** which cannot be `NULL`
  - ◊ Multiple entries may be `NULL` in the same column
  - ◊ Multiple fields can be marked as unique
  - ◊ **Syntax:** `some_field any_type UNIQUE`
  - ◊ Tuples of fields can be marked as unique
  - ◊ **Syntax:** `UNIQUE (field1, fiels2)`
- **CHECK**
  - ◊ Boolean check based on values of a single tuple
  - ◊ Reject if `False`
  - ◊ Accept if `True` or `Unknown`
  - ◊ Some engines treat check somewhat weirdly
  - ◊ Some engines allow subqueries to be part of a check
  - ◊ **Syntax:** `CHECK(some_expression_evaluation_to_bool)`
- **FOREIGN KEY**
  - ◊ Involve two relations
  - ◊ Field must be `NULL` or a valid reference to another table
  - ◊ The reference field is often a **PRIMARY KEY** or at least **UNIQUE**
  - ◊ **Referencing Table:** Table which references a tuple form another table
  - ◊ **Referenced Table:** Table being referenced by another table
  - ◊ **Syntax:** `FOREIGN KEY some_field any_type REFERENCES some_table(some_field)`
  - ◊ **Maintenance**

- \* Changes to the referenced table influences the referencing table
  - And not the other way around!
  - On UPDATE or DELETE
- \* Different ways of handling changes
- \* **Cascade**
  - Propagate modification or delete
- \* **Restrict**
  - Prevent modification or deletion if it violates constraint
    - ▷ By throwing an error
  - Check right after each command
- \* **No Action**
  - Prevent modification or deletion if it violates constraint
    - ▷ By throwing an error
  - Check after a transaction
  - Is the default of PostgreSQL
  - Is equivalent to restrict in MySQL
- \* **Set Default/Set Null:**
  - Set reference to default value or NULL
- \* **Syntax:** ON UPDATE method, ON DELETE method
- Using CONSTRAINT some\_name we can give a name to constraints
  - ◊ Useful in practice since it allows easy modification later on

## 9 Recursive Queries

- Repeatedly execute the same query
- Stop when it converges
  - ◊ I.e. when answer does not change
- Steps
  - ◊ Set R = Empty
  - ◊ Run (base query UNION recursive query) and set it as the new R
    - \* Repeat until R does not change
  - ◊ Query R
- SQL
  - ◊ WITH RECURSIVE some\_table(some\_attribute) AS
    - (<base query>
    - UNION
    - <recursive query>)
    - <Final Query involving some\_table and other relations>
  - \* some\_table(some\_attribute) is the recursive table
    - ◊ Can have any number of attributes
  - \* <base query> is run only once at the beginning and populates some\_table
  - \* <recursive query> queries some\_table and updates it
  - \* <Final Query> run once after some\_table converges
- Recursion can be infinite
  - ◊ Will not terminate
  - ◊ An error will be thrown
- UNION ALL makes bag semantics and may cause infinite recursion when replaces UNION
- Relational model is not well suited for recursion
  - ◊ SQL is based on FOL and FOL cannot express recursion
  - ◊ Other DBs types support recursion more nicely



## 10 NULL

- NULL is a state and not a value
  - ◊ Check: `some_value IS NULL`
    - \* `(NULL IS NULL) -> TRUE`
  - ◊ And not: `some_value = NULL`
    - \* `(NULL = NULL) -> UNKNOWN`
  - ◊ We cannot compare NULL, but we can check if it is NULL
- **Arithmetic:** Always gives NULL if an operand is NULL
- **Comparison:** Always gives UNKNOWN if one of the comparators is NULL
- **Logical operator:** Treats NULL as UNKNOWN
  - ◊ Returns UNKNOWN when the result depends on the concrete value of the variable assigned to NULL
- **Aggregation**
  - ◊ If `group by` a columns containing NULLs, NULL will be in one group
  - ◊ Most aggregation functions ignore the NULL
  - ◊ `Count(*)` ignores NULL not, `Count(column)` ignores it
- Some operators may introduce NULL
  - ◊ E.g. (Left/Right) outer join

## 11 Views

- It is an alias for a query
- Provides higher abstraction than relations
  - ◊ Provides logical data independence
- **Syntax:** `CREATE VIEW some_name AS some_query`
- Can be used similarly as a table
- DMDB convert statement containing a view into a statement without view
  - ◊ I.e. DMDB converts view to a select query
- Used for
  - ◊ **Privacy:** Give person access to a limited view and not the whole relation/DB
  - ◊ **Usability:** Simplify queries
- **Update View**
  - ◊ For base relation  $R_1, \dots, R_n$  and view  $V = Q(R_1, \dots, R_n)$
  - ◊ Update  $V$  into  $V'$  requires finding a set of updated to the base relation  $(R'_1, \dots, R'_n) = f(R_1, \dots, R_n)$  s.t.  $Q(R'_1, \dots, R'_n) = V'$
  - ◊ Not all view can be updated
    - \* Some data is missing
    - \* Primary key is missing
    - \* Update aggregates result
    - \* etc.
  - ◊ SQL tries to avoid indeterminism
  - ◊ SQL view is updatable iff:
    - \* Involved only one base relation
    - \* Involves the key of that base relation
    - \* Does no involve aggregates, group by or duplicate-elimination

## 12 Functional Dependency

- **Redundancy**
  - ◊ Keep same data in multiple relations and/or multiple tuples
    - Waste of storage space
    - Additional work to keep consistent
      - \* Else we get anomalies
    - Hard to keep consistent
    - Additional code to keep consistent
  - + Improve locality
  - + Better performance
  - + Fault tolerance
  - + Availability
- FD is one way to model and understand redundancy
- Models and helps to reason about redundant data
- **FD Definition**
  - ◊ For:
    - \* **Relation Schema:**  $R(A : D_A, B : D_B, C : D_C, D : D_D)$
    - \* **Instance:**  $R \subseteq D_A \times D_B \times D_C \times D_D$
  - ◊ Let  $\alpha \subseteq R, \beta \subseteq R$ 
    - \* Subset of columns
  - ◊ **Functional Dependency**  $\alpha \rightarrow \beta$ : iff  $\forall r, s \in R. r.\alpha = s.\alpha \implies r.\beta = s.\beta$ 
    - \* I.e.  $\alpha \rightarrow \beta \iff$  for any two tuples  $r$  and  $s$  in DB instance  $R$ , if  $r$  and  $s$  share the same values on columns  $\alpha$ , then they share the same values on column  $\beta$
    - \* I.e. there is a mapping, mapping values in columns  $\alpha$  to values in columns in  $\beta$
    - \* **Notation:**  $R \models \alpha \rightarrow \beta$  if  $R$  satisfies  $\alpha \rightarrow \beta$
- **FD  $\alpha \rightarrow \beta$  is minimal:** iff  $\forall A \in \alpha. (\alpha \setminus \{A\}) \not\rightarrow \beta$ 
  - ◊ **Notation:**  $\alpha \rightarrow \cdot \beta$
- **Keys**
  - ◊ **Superkey** is  $\alpha \subseteq \mathcal{R} \iff \alpha \rightarrow \mathcal{R}$ 
    - \* I.e. if we know the value of columns  $\alpha$  we know all values of all columns
    - \* All columns together are a trivial candidate key
  - ◊ **Candidate Key** is  $\alpha \subseteq \mathcal{R} \iff \alpha \rightarrow \cdot \mathcal{R}$ 
    - \* I.e. a minimal super key
- **Implication/Inference**
  - ◊ Given:
    - \* Set  $F$  of some FDs on schema  $\mathcal{R}$
    - \* FD  $\alpha \rightarrow \beta$  which is  $\notin F$
  - ◊  $F$  **implies**  $\alpha \rightarrow \beta$  if every relation instance  $R$  of  $\mathcal{R}$  that satisfies all FDs on  $F$  also satisfies  $\alpha \rightarrow \beta$
  - ◊ **Notation:**  $F \models \alpha \rightarrow \beta$
- **Derivation**
  - ◊ Given:
    - \* Set  $F$  of some FDs on schema  $\mathcal{R}$
    - \* FD  $\alpha \rightarrow \beta$  which is  $\notin F$
  - ◊  $F$  **derives**  $\alpha \rightarrow \beta$  if there is a derivation (using Armstrong's axioms) from  $F$  to  $\alpha \rightarrow \beta$
  - ◊ **Notation:**  $F \vdash \alpha \rightarrow \beta$
- **Closure**

- ◇ Given:
  - \* Set  $F$  of some FDs on schema  $\mathcal{R}$
  - \* Set of attributes  $\alpha \subseteq \mathcal{R}$
- ◇ **Closure** of  $\alpha$  is the set of all attributes  $y \in \mathcal{R}$ , such that  $\alpha \rightarrow y$  can be derived from  $F$
- ◇ **Notation:**  $\alpha^+$
- ◇  $\alpha^+ = \{y \in \mathcal{R} \mid F \vdash \alpha \rightarrow y\}$
- ◇  $F \vdash \alpha \rightarrow \beta \iff \beta \subseteq \alpha^+$
- ◇ **Recipe:** Find  $\alpha^+$ 
  - 1) Set  $\alpha^+$  to  $\alpha$
  - 2) For all  $\beta \in \alpha$  check if there is a  $\beta \rightarrow \gamma$  and add  $\gamma$  to  $\alpha^+$  if so
  - 3) Repeat step 2 until  $\alpha^+$  converges
- **Minimal Basis/Cover**
  - ◇ **Goal:** Remove redundant FDs in a set of FDs
    - \* I.e. FDs which we can derive from the other FDs
  - ◇ Given: Set  $F$  of some FDs
  - ◇ A **Minimal Cover** of  $F$  is a set  $G \subseteq F$  with:
    - \*  $G \equiv F$
    - \* All FDs in  $G$  have the form  $X \rightarrow \alpha$ , where  $\alpha$  is a single attribute
    - \* It is not possible to make  $G$  smaller by:
      - $G \setminus \{X \rightarrow \alpha\} \not\equiv G, \forall X \rightarrow \alpha \in G$   
▷ I.e. remove a FD
      - $(G \setminus \{X \rightarrow \beta\}) \cup \{X \rightarrow \beta\} \not\equiv G \forall X \rightarrow \beta \in G$   
▷ I.e. Remove an attribute from a FD
  - ◇ **Recipe:** Compute minimal basis
    - 1) Set  $G$  to the set of FDs obtained from  $F$  when decomposing the RHS of each FD to a single attribute
    - 2) Remove trivial FDs
      - I.e. all  $a \rightarrow b$  where  $b \subseteq a$
    - 3) Remove all redundant attributes from the LHS of all FDs
      - I.e. if  $a \rightarrow b$  and  $\exists x \in a$  such that  $(a \setminus x) \rightarrow b$ , replace  $a \rightarrow b$  with  $(a \setminus x) \rightarrow b$
    - 4) Remove all redundant FDs
      - I.e. if  $a \rightarrow b$  and  $\{b\} \subseteq \text{Closure}(F \setminus \{a \rightarrow b\}, a)$ , remove  $a \rightarrow b$  from  $F$
- **Equivalence**
  - ◇ Given: Set  $F$  and  $G$  of FDs on schema  $\mathcal{R}$
  - ◇  $F$  and  $G$  are **equivalent** if  $F \models G$  and  $G \models F$
  - ◇ **Notation:**  $F \equiv G$
  - ◇ Cardinalities define FD
  - ◇ FD determine keys
- **Armstrong Axioms**
  - ◇ Axioms:
    - \* **Reflexivity:**  $\alpha \subseteq \beta \implies \beta \rightarrow \alpha$ 
      - Special case is  $\mathcal{R} \rightarrow \alpha$
      - Trivial FDs
    - \* **Augmentation:**  $\alpha \rightarrow \beta \implies \alpha\gamma \rightarrow \beta\gamma$ 
      - Where  $\alpha\gamma := \alpha \cup \gamma$
    - \* **Transitivity:**  $\alpha \rightarrow \beta \wedge \beta \rightarrow \gamma \implies \alpha \rightarrow \gamma$
  - ◇ These axioms are both:
    - \* **Sound:**  $F \vdash \alpha \rightarrow \beta \implies F \models \alpha \rightarrow \beta$

- \* **Complete:**  $F \models \alpha \rightarrow \beta \implies F \vdash \alpha \rightarrow \beta$
- ◇ All other FDs can be implied from these axioms
- **Other Rules**
  - ◇ **Union:**  $\alpha \rightarrow \beta \wedge \alpha \rightarrow \gamma \implies \alpha \rightarrow \beta\gamma$
  - ◇ **Composition:**  $\alpha \rightarrow \beta\gamma \implies \alpha \rightarrow \beta \wedge \alpha \rightarrow \gamma$
  - ◇ **Pseudo Transitivity:**  $\alpha \rightarrow \beta \wedge \beta\gamma \rightarrow \theta \implies \alpha\gamma \rightarrow \theta$
- **Composition of Relations**
  - ◇ **Goal:** Split bad relations into ones containing only one concept
    - \* **Bad Relation:** Relation which combine several concepts
  - ◇ **Lossless Decomposition** Of  $R$  into  $R_1, \dots, R_n$  if  $R = R_1 \bowtie R_2 \bowtie \dots \bowtie R_n$ 
    - \* For  $\mathcal{R} = \mathcal{R}_1 \cup \mathcal{R}_2$ , decomposition  $R_1 = \Pi_{\mathcal{R}_1}(R), R_2 = \Pi_{\mathcal{R}_2}(R)$  is lossless, if  $(\mathcal{R}_1 \cap \mathcal{R}_2) \rightarrow \mathcal{R}_1 \vee (\mathcal{R}_1 \cap \mathcal{R}_2) \rightarrow \mathcal{R}_2$
  - ◇ **Recipe:** Show given decomposition  $R_1, R_2$  is lossy
    - \* Decomposition is lossless iff  $R = R_1 \bowtie R_2$
    - \* Find some instance of the relations which serves as a counterexample
  - ◇ **Recipe:** Show given decomposition  $R_1, R_2$  is lossless
    - \* Decomposition is lossless if  $(R_1 \cap R_2) \rightarrow R_i, i \in \{1, 2\}$
  - ◇ **Preservation of Dependencies:**  $\text{FD}(R)^+ = [\text{FD}(R_1) \cup \dots \cup \text{FD}(R_n)]^+$

## 13 Normal Form

- **Normalisation:** Process of restructuring a DB to reduce data redundancy and improve data integrity
  - ◊ Done using a series of normal forms
- **Normal Forms:** Common way for normalization
- **First Normal Form (1NF)**
  - ◊ All values are of atomic domains
    - \* I.e. only int, double, char etc. and not tuples, lists etc.
- **Second Normal Form (2NF)**
  - ◊  $R$  is in 2NF iff every non-key attribute is minimally dependent on every key
    - \* **Minimally Dependant:** No attribute depends on part of a key
    - \* I.e. none of the non-key attribute depends on part of a key
  - ◊ I.e. If  $R(a, b, c, d)$  with primary key  $\{a, b\}$  then  $a \rightarrow c$  is not allowed
  - ◊ Alternative, less strong definition
  - ◊  $R$  is in 2NF if for all  $\alpha \rightarrow B$  at least one holds:
    - \*  $B \in \alpha$
    - \*  $B$  is an attribute of at least one key
    - \*  $\alpha$  is a superkey of  $R$
    - \* No attribute in  $\alpha$  is part of any key
  - + Improve insert, update and delete anomaly
  - Does not solve update and delete anomaly
    - \* Because we can have  $C \rightarrow D$  where both are non-keys
  - ◊ Enforce 2NF by decomposing the relation into multiple relations
- **Third Normal Form (3NF)**
  - ◊  $R$  is in 3NF iff for all  $\alpha \rightarrow B$  at least one holds:
    - \*  $B \in \alpha$
    - \*  $B$  is an attribute of at least one key
    - \*  $\alpha$  is a superkey of  $R$
  - ◊ I.e. If  $\alpha \rightarrow \beta$  does not satisfy any of these conditions,  $\alpha$  is a concept on its own
    - \* Gets rid of transitive dependency
  - Does not get rid of all redundant data
  - + Is lossless
  - + Preserves all dependencies
  - ◊ **Recipe Synthesis Algorithm:** Decompose  $\mathcal{R}$  into  $\mathcal{R}_1, \dots, \mathcal{R}_n$  according to 3NF
    - 1) Compute minimal basis  $F_c$  of  $F$
    - 2) For all  $\alpha \rightarrow \beta \in F_c$ , create  $R_{\alpha \cup \beta}(\alpha \cup \beta)$
    - 3) If none of the relations contains a superkey, add a relation with a key
    - 4) Eliminate  $R_\alpha$  if there exists  $R_{\alpha'}$  such that  $\alpha \subseteq \alpha'$
- **Boyce-Cod Normal Form (BCNF)**
  - ◊  $R$  is in BCNF iff for all  $\alpha \rightarrow \beta$  at least one holds:
    - \*  $B \in \alpha$
    - \*  $\alpha$  is a superkey of  $R$
  - ◊ I.e. Each relation stores the same information only once
  - Does not preserve all FDs
  - Does not get rid of all data redundancies
    - \* Only of all redundancies caused by FD
  - ◊ **Recipe Decomposition Algorithm:** Decompose  $\mathcal{R}$  into  $\mathcal{R}_1, \dots, \mathcal{R}_n$  according to BCNF

- 1) Set result to  $\{\mathcal{R}\}$
  - 2) If there is  $\mathcal{R}_i$  with  $\alpha \rightarrow \beta$  which is not in BCNF
    - $\mathcal{R}_i^1 = \alpha \cup \beta$
    - $\mathcal{R}_i^2 = \mathcal{R}_i \setminus \beta$
    - $\text{Result} = (\text{Result} \setminus \mathcal{R}_i) \cup \{\mathcal{R}_i^1, \mathcal{R}_i^2\}$
  - 3) Repeat 2 as long as there are  $\mathcal{R}_i$  which are not in BCNF
- **Non-First Normal Form (NFNF)**
    - ◊ Use an array to get rid of data redundancy
      - \* Will not be in 1NF
    - ◊ Can be used in SQL
  - **4th Normal Form**
    - ◊ **Multi-Value Dependency (MVD)**
      - \*  $A$  is MVD on  $B$  and  $C$  means that the value of  $B$  does not have impact on the value of  $C$ , and that  $B$  and  $C$  can take multiple values for the same  $A$ .
      - \* **Notation:**  $A \twoheadrightarrow B, A \twoheadrightarrow C$
      - \*  $\alpha \twoheadrightarrow \beta$  for  $R(\alpha, \beta, \gamma)$  iff
        - $\forall t_1, t_2 \in R, t_1.\alpha = t_2.\alpha \implies \exists t_3, t_4 \in R:$ 
          - ▷  $t_3.\alpha = t_4.\alpha = t_1.\alpha = t_2.\alpha$
          - ▷  $t_3.\beta = t_1.\beta; t_4.\beta = t_2.\beta$
          - ▷  $t_3.\gamma = t_2.\gamma; t_4.\gamma = t_1.\gamma$
      - \* **Intuitively:** Thinks about in terms of joins
        - $R(\alpha, \beta, \gamma)$  with  $\alpha \twoheadrightarrow \beta$  can be decomposed into  $R = R_1 \bowtie R_2$ 
          - ▷  $R_1 = \Pi_{\alpha, \beta} R$
          - ▷  $R_2 = \Pi_{\alpha, \gamma} R$
          - ▷ Is lossless if  $\alpha \twoheadrightarrow \beta$  or  $\alpha \twoheadrightarrow \gamma$
      - \* Can result in anomalies and redundancy
      - \* **Trivial:**
        - $\mathcal{R}(\alpha, \theta) : \alpha \twoheadrightarrow \alpha\theta$ 
          - ▷ I.e.  $\alpha \twoheadrightarrow \mathcal{R}$
        - $\mathcal{R}(\alpha, \theta) : \alpha \twoheadrightarrow \theta$ 
          - ▷ I.e.  $\alpha \twoheadrightarrow (\mathcal{R} \setminus \alpha)$
        - $\beta \subseteq \alpha \implies \alpha \twoheadrightarrow \beta$
      - \* **Promotion:**  $\alpha \rightarrow \beta \implies \alpha \twoheadrightarrow \beta$
      - \* **Complement:**  $\alpha \twoheadrightarrow \beta \implies \alpha \twoheadrightarrow (\mathcal{R} \setminus \alpha \setminus \beta)$
      - \* **Multi-Value Augmentation:**  $\alpha \twoheadrightarrow \beta \wedge (\delta \subseteq \gamma) \implies \alpha\gamma \twoheadrightarrow \beta\delta$
      - \* **Multi-Value Transitivity:**  $(\alpha \twoheadrightarrow \beta) \wedge (\beta \twoheadrightarrow \gamma) \implies \alpha \twoheadrightarrow \gamma$
      - \* Not all FD rules apply to MVD
        - Need to distinct between FD and MVD
    - ◊ Deals with MVD (and not FD)
    - ◊  $R$  is 4NF iff for all  $\alpha \twoheadrightarrow \beta$ , at least one condition holds:
      - \*  $\alpha \twoheadrightarrow \beta$  is trivial
      - \*  $\alpha$  is a superkey of  $R$
    - ◊  $R$  in 4NF  $\implies R$  in BCNF
    - ◊ **Recipe Decomposition Algorithm:** Decompose  $\mathcal{R}$  into  $\mathcal{R}_1, \dots, \mathcal{R}_n$  according to 4NF
      - 1) Set result to  $\{\mathcal{R}\}$
      - 2) If there is  $\mathcal{R}_i$  with  $\alpha \twoheadrightarrow \beta$  which is not in 4NF
        - $\mathcal{R}_i^1 = \alpha \cup \beta$
        - $\mathcal{R}_i^2 = \mathcal{R}_i \setminus \beta$

- $\text{Result} = (\text{Result} \setminus R_i) \cup \{\mathcal{R}_i^1, \mathcal{R}_i^2\}$
  - 3) Repeat 2 as long as there are  $\mathcal{R}_i$  which are not in 4NF
- lossless
- preserve dependencies
- $\underbrace{1\text{NF} \subset 2\text{NF} \subset 3\text{NF} \subset \text{BCNF}}_{\text{deal with FD}} \subset \underbrace{4\text{NF}}_{\text{deal with MVD}}$
  - There are many different NF with different properties
  - **Denormalisation**
    - ◊ Higher normalisation is not always better
    - ◊ Sometimes we deliberately denormalize a DB
    - + Faster due to better locality
    - More redundant data



## 14 Analytic

- **This section is not exam relevant! TODO: Finish summarizing this section**
- Look at DB from a statistical view
- Useful for
  - ◊ Data mining
  - ◊ Machine Learning

### 14.1 Associated Rule Mining

- Aka Data Mining
- **Frequent Itemsets I**: Items appearing at least in  $s$  transaction together
  - ◊  $s$ : Support threshold
- **Support**: Number of transactions containing all items of  $I$
- **Association Rules**
  - ◊ We can say what item is likely in a transaction knowing parts of the transaction
  - ◊  $\{i_1, \dots, i_k\} \rightarrow j$ : If  $i_1, \dots, i_k$  in transaction, then  $j$  in transaction
  - ◊ Allows creating suggestions/recommendations
  - ◊ There are many more such rules
    - \* Only interesting in relevant ones
  - ◊ **Confidence  $c$** : Probability that transaction contains  $j$  given  $i_1, \dots, i_k$ 
    - \*  $\text{conf}(I \rightarrow j) = \frac{\text{support}(I \cup j)}{\text{support}(I)}$
    - \* Number of transaction with  $i_1, \dots, i_k, j$  divided by the number of transaction with  $i_1, \dots, i_k$
  - ◊ Not all high-confidence rules are interesting
    - \* An item may be in many transaction anyways
  - ◊ **Interest**: Difference between its confidence and the fraction of baskets that contain  $j$ 
    - \*  $|\text{confidence} - \# \text{ baskets with } j|$
    - \* Rules are interesting if value is  $\geq \sim 5$
- **Mining Association Rules**
  - ◊ Find all frequent itemsets  $I$ 
    - \* **Naive Algorithm**:
      - Brute force
      - Each itemset is a candidate
      - Time:  $\mathcal{O}NMw$
      - Space:  $\mathcal{O}M$
      - ▷  $M = 2^d$
    - \* **A-Priori**
      - **Idea**:
        - ▷ If itemset is frequent, then all of its subsets must also be frequent
        - ▷ If itemset is not frequent, then no superset will be frequent
        - ▷ Support of itemset never exceeds the support of its subsets
      - $C_k$ : candidate itemset of size  $k$
      - $L_k$ : frequent itemset of size  $k$
      - Steps
        - ▷ **Initial**:  $k = 1, C_1 = \text{all items}$
        - ▷ While  $C_k$  is not empty
          - Scan DB to find which itemsets in  $C_k$  are frequent and put them into  $L_k$

- Use  $L_k$  to generate a collection of candidate itemsets  $C_{k+1}$  of size  $k + 1$ 
  - Join two itemsets of size  $k$  that share the first  $k - 1$  items
  - Use principles to filter
  - $k = k + 1$
- There are libraries for that
- ◊ For every subset  $A$  of  $I$ , generate a rule  $A \rightarrow I/A$
- ◊ Output rules above the confidence threshold

## 14.2 Clustering

- Group objects into different categories
- Each cluster is a subset of **TODO: ?**
- **Clustering:** Set of clusters
- **Partition of Clustering:**
  - ◊ A division date objects into subsets (clusters) such that each data object is in exactly one subset
- **Hierarchical Grouping:**
  - ◊ Tree
- **Traditional Hierarchical Clustering:** Each is a subset of each other
- **Non-Traditional Hierarchical Clustering:** Different groups
- Types
- **Well-Separated Clusters:**
  - ◊
- **Center-Based:**
  - ◊ Elements are closest to center of their own cluster
- **Contiguity-Base:**
  - ◊ Set of points such that a point in a cluster is closer to one or more other points in the cluster than to any points not in the cluster
- **Density-Based:**
  - ◊
- **Conceptual Clusters:**
  - ◊ Cluster of objects which share, or not share a certain object
- **K-Means**
  - ◊ Partitional clustering approach
  - ◊ Each cluster is associated with a centroid
  - ◊ Each point is assigned to the cluster with the closest centroid
  - ◊ Number of clusters  $K$  is given
  - ◊
  - ◊ Often uses Euclidean distance
  - ◊ Minimizing the Sum of Square Error (SSE)
  - ◊ NP-Hard if  $d \geq 2$
  - ◊ Polynomial if  $d = 1$
  - ◊ Can be estimated
  - ◊ Iterative algorithm

Randomly select  $K$  point is the initial centroids

repeat

From  $K$  clusters by assigning all points to the closest centroid

Recompute

  - ◊ Initial starting chaise results in different results

- \* Do multiple runs and choose best result
- ◇ Centroid depends on distance function
  - \* NP-hard to find centroid for some functions
- + Will always converge
- + Quick convergence
- ◇  $\mathcal{O}nKI d$ 
  - \* **n**: number of points
  - \* **I**: number of iterations
  - \* **K**: number of clusters
  - \* **d**: dimension
- Problematic under certain conditions
- ◇ Better control if start with many clusters and merge then later on into fewer
- ◇ MADlib for SQL support

### 14.3 Classification

- Different from clustering
- Learning a target function  $f$  that maps attribute set  $x$  to one of the predefined class labels  $y$
- Training set consists of records with known class labels
- Training set is used to build classification model
- 
- 
- **Instance-Based Learning:**
  - ◇
- **Nearest Neighbour Classifier:**
  - ◇ Requirements
    - \* Set of stored records
    - \* Distance metrics to compute distance between records
    - \* The value of  $k$ 
      - Number of nearest neighbours

## 15 DMDB System

- Complex application
- Simplification we consider
  - ◊ All data are stored on disk
  - ◊ Disk is larger than memory
  - ◊ Disk is slower than memory
  - ◊ Disk favours different access pattern
  - ◊ Single CPU
  - ◊ One relation is stored in a single file
- Overview
  - ◊ **Query Optimization:** SQL query to relational algebra converter
  - ◊ **Operator Execution:** Execute a relation algebra operation
  - ◊ **Access Methods:** Provides different ways of accessing data from a relation
  - ◊ **Buffer Pool Management:** Gives illusion that all data is stored on memory

### 15.1 Storage Hierarchy

- Storage is a hierarchy
- Challenge: keep CPU busy
- Rapidly changing
- Different hierarchies lead to different DB design
- **Hard Drive**
  - ◊ Plates spin
  - ◊ Plate is split into sectors of fixed size
  - ◊ Arm assembly is moved in or out to position a head on a desired track
  - ◊ Tracks under head makes a cylinder (kind off)
  - ◊ Only one head reads/writes at any time
  - ◊ Block size is a multiple of sector size
  - ◊ Performance
    - \* **Seek time  $t_s$ :** Moving arm to position disk head on track
      - 10 – 20ms
    - \* **Rotate time  $t_r$ :** waiting for block to rotate under head
      - 10 – 20ms
    - \* **Transfer time  $t_{tr}$ :** actually moving data to/from disk surface
      - 8KB/0.1ms
    - \* **Random access D times:**  $D(t_s + t_r + t_{tr})$
    - \* **Sequential access D time:**  $t_s + t_r + Dt_{tr}$
- We consider abstraction HD → DRAM → CPU
- DRAM → CPU is much faster than HD → DRAM
  - ◊ We only consider HD → DRAM optimisation

### 15.2 Disk Manager

- Lowest level
- Used by higher levels for
  - ◊ Allocate/de-allocate pages
  - ◊ Read/Write pages
- Requests for sequential allocation must be satisfied
- Responsible for maintaining a database's files

- DB content is stored as one or multiple files
- Many DMDBs use the file system provided by the OS
  - ◊ People used to build custom file systems for DMDBs
- **Files** contain a collection of pages
- **Page** is a fixed-size block of data
  - ◊ Contains a collection of tuples
  - ◊ **page id**: unique identifier of each page
- A relation is stored as a collection of pages

### 15.2.1 File Layout

- How does a file manage all its pages?
- Unordered collection of pages
- Support record level operations
- We must keep track of:
  - ◊ the pages in the file
  - ◊ the records on each page
  - ◊ free space on each page
- **Heap File**:
  - ◊ Unordered collection of pages
  - ◊ Need to keep track where tuples are stored and of free space
  - ◊ Two ways to implement it
  - ◊ **Linked List**
    - \* **Header Page**: One single page
      - Kind of the root
      - Has to pointer to two linked lists
        - ▷ Free pages list
        - ▷ Data pages list
    - No global view on data
    - \* Performance
      - Assume
        - ▷ Directory fits in and is in memory
        - ▷  $\#Pages = D$
        - ▷ Pages are randomly allocated on disk
          - Models worst-case
      - **Insert**:  $t_{s+r} + 2t_{trans}$ 
        - ▷ If page 1 has slot available
      - **Find Record**: By non-RID value (RID is a pair page, id and slot id)
        - ▷  $\frac{D}{2}(t_{r+s} + t_{trans})$
      - **Scan**:
        - ▷  $D(t_{s+r} + t_{trans})$
- ◊ **Page Directory**
  - \* **Header Page**: Multiple pages
    - Each contains a list of pointer to data pages
  - \* Performance
    - Assume
      - ▷ Directory fits in and is in memory
      - ▷  $\#Pages = D$
      - ▷ Pages are sequentially allocated on disk
        - Models best case

- **Insert:**  $t_{s+r} + 2t_{trans}$
- **Find Record:** By non-RID value (RID is a pair page, id and slot id)
  - ▷  $t_{s+r} + \frac{D}{2}t_{trans}$
- **Scan:**
  - ▷  $t_{s+r} + Dt_{trans}$

### 15.2.2 Page Layout

- Page consists of
  - ◊ **Header:** Contains metadata like
    - \* Page size
    - \* DBMS version
    - \* Compression information
    - \* Encryption information
    - \* Checksum
  - ◊ **Data:** Actual tuples
- Header is at top of page
- Data starts after header
- Multiple strategies
- **Naive Strategy**
  - ◊ Page is split into slots of fixed size
    - \* One tuple goes into one slot
  - ◊ Header keeps track of occupied/free slots
    - Lots of space wasted when tuples are of different length
- **Slotted Page**
  - ◊ Record id =  $\langle \text{page id, slot \#} \rangle$
  - ◊ **Slot array:** Array of pointers and size of occupied slots
    - \* Comes right after the header
  - + Can move tuples on page without changing record id

### 15.2.3 Tuple layout

- Data access methods:
- **On-line transaction processing (OLTP):**
  - ◊ Simple query
  - ◊ Reads/writes a small amount of data related to a single entry
- **On-line analytical processing (OLAP):**
  - ◊ Complex queries
  - ◊ Read large portions of the DB spanning multiple entries
- Multiple implementations
- **Row Storage**
  - ◊ Store tuple together
  - ◊ Divided into
    - \* **Bitmap:** Indicates which attributes are NULL (somehow)
    - \* **Fixed-Length:**
      - ◊ Contains fixed-length fields
      - ◊ Arrangement and sizes are equal for all tuples
      - ◊ Can directly access the  $i$ -th field
    - \* **Variable-Length:**
      - ◊ Contains variable-length fields

- Two implementation
  - **Field Delimited:** Special characters mark the end/start of fields
    - Access  $i$ -th fields requires scann of list
  - **Field Offset Array:** Array where  $A[i]$  contains the start of the  $i$ -th field
    - ▷ Stored at the beginning of the tuple
    - + Direct access to  $i$ -th field
- + All tuple information are together
- + Good for OLTP
  - Bad for OLAP
    - \* Read lots of data we do not care
- **Column Storage**
  - ◊ Store a whole column together
  - + Good for OLAP
    - Slow for point queries (look for a single value), inserts, updates and deletes
    - Bad for OLTP
  - + Easier data compression

### 15.3 Buffer Pool Management

- Buffer manager acts like the intermediate layer between the system and disk manager
- On fetch, check if desired page is in RAM. If not, bring to RAM and returns. Else directly return it.
- **Goal:** Provide illusion that all data in in RAM
- Page Replacement
  - ◊ If RAM is full, we need to evict a page
  - ◊ Eviction policy is of key importance
  - ◊ **Future access pattern known**
    - \* **Idea:** Evict block whose next access is farthest in the future
    - \* Called Belady's MIN algorithm
    - + Optimal under this assumption
  - ◊ **Future access pattern unknown**
    - \* Rarely know about the future access pattern
    - \* Different strategies
    - \* **Least Recently Used (LRU):** Evict the least recently used page
      - + Works well for repeated access to popular pages
      - 100% miss under sequential flooding
        - ▷ **Sequential Flooding:** Access in a repeated pattern but such that not all pages fit into RAM
      - At most twice as bad compared to optimal when LRU has twice the memory  
**TODO: What does this mean?**
    - \* **Most Recently Used (MRU):** Evict the most recently used page
      - Frequently accessed page has to be fetched often
      - ▷ E.g. Index scan
    - \* Access patterns
      - **Sequential:** Table Scan:  $1 \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow \dots$
      - **Hierarchical:** Index Scan:  $1 \rightarrow 4 \rightarrow 11 \rightarrow 1 \rightarrow 4 \rightarrow 12 \rightarrow 1 \rightarrow 3 \rightarrow 8 \dots$
      - **Random:** Index Lookup:  $12 \rightarrow 9 \rightarrow 4 \rightarrow 21 \rightarrow 55 \rightarrow 6 \rightarrow 42 \rightarrow \dots$
      - **Cyclic:** Nested-Loop:  $1 \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow 1 \rightarrow 2 \rightarrow 3 \rightarrow 4 \rightarrow \dots$
- \* Depending on the pattern different strategies perform differently well
- \* If we have some information about pattern, we can select an optimal policy

- This is the key different from the OS RAM manager
- OS vs DB
  - ◊ Principles are similar, but in DB we know more what is going on
  - ◊ Allows for better optimisation
  - ◊ DB uses own buffer
  - ◊ DB uses OS's filesystem
    - \* Has own buffer which can cause problems

## 15.4 Access Methods

- Used by upper layer to access information in tables
- Provides different way of accessing data from a relation
- Goal: find a tuple whose attribute  $X$  equals  $Y$
- **Sequential Scan:**
  - ◊ Bring page 1 and scan, bring page 2 and scan etc.
  - ◊ Cost:
    - \* When pages are sequentially allocated:  $t_{s+r} + D(t_{tr})$
    - \* When pages are randomly allocated:  $D(t_{s+r} + t_{tr})$
    - \* When write back on scan pay extra:  $t_{tr}$
- Improve: build data structure  $f$  over relation, which can easily answers our scan
  - ◊ Evaluation of  $f$  is usually cheaper than sequential scan
- **B-Tree (B+ Tree)**
  - ◊ Self-balancing tree that keeps data sorted
  - ◊ Properties
    - \*  $M$ -way search tree
      - Like a binary search tree but each node has up to  $M$  children
    - \* Perfectly balanced
    - \* Every inner node (except root) is at least half-full
    - \* Every inner node with  $k$  keys has  $k + 1$  non-null children
    - \* Each page is represented as a node
    - \*  $\mathcal{O}(\log n)$  for search, insertion, deletion
  - ◊ Two flavours
    - \* **Unclustered B+ Tree:** Leaf node contains RID (PageID, SlotID), pointing to the relation
    - \* **Clustered B+ Tree:** Leaf node contains the actual tuple
      - Can have only one clustered B+ Tree per relation
  - ◊ **Point query**
    - \* Find single key
    - \* Node size =  $M$ , # Tuples =  $N$
    - \* Depth  $\mathcal{O}(\log_M N)$
    - \* Num I/O
      - Unclustered:  $\underbrace{\log_M N}_{\text{search tree}} + \underbrace{1}_{\text{access actual tuple}}$
      - Clustered:  $\underbrace{\log_M n}_{\text{search tree}}$
    - \* Much faster than sequential scan
  - ◊ **Range query**
    - \* Find all tuples in range
    - \* Node size =  $M$ , # Tuples =  $N$
    - \* Depth  $\mathcal{O}(\log_M N)$



- \* Num I/O

- Unclustered:  $\underbrace{\log_M N}_{\text{search tree}} + \underbrace{\frac{\# \text{tuples}}{\text{tuples-per-page1}}}_{\# \text{ leafs to read to get all RIDs}} + \underbrace{\# \text{tuples}}_{\text{cost of reading}}$

- Clustered:  $\underbrace{\log_M N}_{\text{search tree}} + \underbrace{\frac{\# \text{tuples}}{\text{tuple-per-page2}}}_{\# \text{ leafs to read to bet all RIDs}}$

- tuple-per-page2 < tuple-per-page1 since clustered contains the actual tuples (which are larger than RID)

- ◇ **Insert**

- \* Insert a tuple

- \* Algorithm

- Find the right leaf node  $L$
    - Put data in  $L$ 
      - ▷ If  $L$  has enough space we are done
      - ▷ Otherwise, split  $L$ , insert key to the parent of  $L$

- ◇ **Delete**

- \* Delete a tuple

- \* Algorithm

- Find the right leaf node  $L$
    - Remove data in  $L$ 
      - ▷ If  $L$  is at least half full we are done
      - ▷ Otherwise, merge two leaf nodes or borrow one tuple from neighbours, update parents

- ◇ Heap file vs B+ Tree

- \* **Heap file:**

- + Lot of sequential scans

- \* **B+ Tree:**

- + Small number of random access

- \* Tradeoff, between few expensive, or many cheap accesses

- ◇ Use Btree: `CREATE INDEX name ON table USING btree(column);`

- ◇ Bulk build

- \* Different approaches

- \* Insert tuple by tuple

- Slow

- \* Sort and insert tuple bottom-up

- + Query time is independent of data distribution

- **Hash Table**

- ◇ Goal: Do better than B+ Trees for point queries

- ◇ **Hash Function  $f(x)$ :** Maps each object into an entry in the table

- \* E.g. For integer:  $h(x) = (ax + b) \bmod p$
    - \* E.g. For Strings:  $h(s_1, \dots, s_n) = (\sum_i s_i a^i) \bmod p$
    - \* Hard to find ideal hash function

- ◇ Ideally  $a \neq b \implies f(a) \neq f(b)$

- \* Else we get a collision

- ◇ Different approaches to prevent collisions

- ◇ **Closed Hashing:** We know how many elements are trying to index

- \* **Linear Probe**

- Working

- ▷ Compute hash to find right slot
  - ▷ Insert object into next empty slot
    - Deletion is non-trivial
      - We have to insert special marker into delete position to mark that this is a contiguous block
  - Search/Insert:  $\mathcal{O}(\text{size of largest cluster})$ 
    - ▷ **Cluster:** Largest consecutive sequence of occupied slots
  - For hash table
    - ▷ Size  $m$
    - ▷ Contains  $n = \lambda m$  keys
    - ▷ Full to  $\lambda\%$
    - we have on average:
      - ▷ **Insert:**  $\frac{1}{2}(1 + \frac{1}{(1-\lambda)^2})$
      - ▷ **(Successful) search:**  $\frac{1}{2}(1 + \frac{1}{(1-\lambda)})$ 
        - Good if  $\lambda \approx 50\%$  full
  - + Very cache-efficient
    - Very sensitive to hash function
      - ▷ Hash function must be “good”
- ◇ **Open Hashing:** We do not know how many elements are there
  - \* **Chained Hashing**
    - Working
      - ▷ Compute hash to find right slot
      - ▷ Insert if no collision
      - ▷ Else append to linked list
    - Expected length of chain for
      - ▷  $h(x)$  is uniformly random
      - ▷  $m = \mathcal{O}(N)$
      - ▷  $N$ : Number of slots

is  $\mathcal{O}(1)$

## 15.5 Operator Execution

- Execute a relational algebra operator
- Used different access methods to implement these operators
- **Select  $\sigma_C(R)$ :**
  - ◇ Input  $R$ , condition  $C$
  - ◇ Assume
    - \*  $R$  has  $|R|$  tuples and  $B(R)$  pages
    - \* Buffer size:  $M$  pages
  - ◇ **Selectivity:**  $\alpha(C, R)$ 
    - \* Is a constant
    - \* Number of tuples in  $R$  that satisfy condition  $C$  divided by  $|R|$
  - ◇ **Sequential Scan (Clustered) Heap File**
    - \* Bring pages to RAM one by one
    - \* Scan each tuple and check for the predicate  $C$
    - \* If true, output tuple
    - \* Total cost:  $\underbrace{\mathcal{O}(B(R))}_{\text{read}} + \underbrace{\alpha(C, R)B(R)}_{\text{write}}$
    - \* Cost is different for each query

- ◊ **Index Scan**
  - \* If  $C \in \{<, >, =\}$
  - \* **Unclustered B+ Tree**
    - Steps
      - ▷ Find the right leaf node
      - ▷ Scan the index
      - ▷ Fetch and return corresponding tuple from heap file
    - Total Cost:  $\mathcal{O}(\underbrace{\log |R|}_{\text{find leaf}} + \underbrace{\alpha(C, R)|R|}_{\text{fetch tuple}} + \underbrace{\alpha(C, R)B(R)}_{\text{write}})$
  - \* **Clustered B+ Tree**
    - Steps
      - ▷ Find the right leaf node
      - ▷ Scan the index
      - ▷ Return tuple
    - Total Cost:  $\mathcal{O}(\underbrace{\log |R|}_{\text{find leaf}} + \underbrace{\alpha(C, R)B(R)}_{\text{fetch tuple}} + \underbrace{\alpha(C, R)B(R)}_{\text{write}})$
- **Sort**
  - ◊ **Sort(R, Attribute A)**
  - ◊ **Clustered B+ Tree**
    - \* Sorted leaves and constant access time
  - ◊ **Unclustered B+ Tree**
    - \* Sorted leaves but one random access per tuple
  - ◊ Different sorting algorithms
  - ◊ Assume: Data  $N$  is much larger than buffer  $B$
  - ◊ **External Sort**
    - \* If  $B = 3$  we can sort 2 pages
    - \* Working
      - If  $N < B$ 
        - ▷ Run quicksort
      - Otherwise
        - ▷ Phase 1: Sorting
          - Partition file into smaller chunks that fit into memory
          - Sort smaller junks
        - ▷ Phase 2: Merging
          - Combine smaller, sorted chunks into large file
    - \* **TODO: Add Cost**
- **Join  $S \bowtie_{\theta} S$** 
  - ◊ **Nested Loop Join**
    - \* foreach tuple  $r$  in  $R$ :
      - foreach tuple  $s$  in  $S$ :
        - if  $\text{Theta}(r, s)$ :
          - output  $r, s$
    - \*  $M = 3$ :  $\mathcal{O}(B(R) + |R|B(S) + \text{print IO})$
    - \* Order of tables matters
      - Depends on table size, buffer size
      - Small  $M$ : Better if smaller relation is in outer loop
      - Large  $M$  (assume smaller table is fully cached): Better if smaller relation is in the inner loop
    - \* **Efficient if:** Both relations fit into memory

◇ **Block Nested Loop Join**

- \* foreach block BR in R:
  - foreach block BS in S:
    - foreach tuple r in BR:
      - foreach tuple s in BS:
        - if  $\Theta(r, s)$ :
          - output r, s
- \*  $M = 3$ :  $\mathcal{O}(B(R) + B(R)B(S))$
- \*  $M > 3$ :  $\mathcal{O}(B(S) + B(R)\frac{B(S)}{M-2})$
- \* Partition  $S$  into  $a = \frac{B(S)}{M-2}$  junks
  - Fully cache junk
  - Pay  $B(R)$  to join this junk with  $R$ :  $(M-2) + B(R)$
  - Repeat  $\frac{B(S)}{M-2}$  times **TODO: Not sure what this means**

◇ **Index Nested Loop Join**

- \* foreach tuple r in R:
  - foreach tuple s in IndexScan( $S, r, \Theta$ ):
    - output r, s
- \*  $\mathcal{O}(B(R) + |R|C)$ 
  - $C$  is the cost of lookup in the index

◇ **Sort Merge Join**

- \* Assume both relations are sorted
- \* Working
  - Scan both relation
  - Compare to head
- \* Cost  $\mathcal{O}(B(R) + B(S) + \text{Sort}(R) + \text{Sort}(S))$
- \* **Efficient if:** Index in attribute is present

◇ **Hash Join**

- \* build Hash Table HT for R
  - foreach tuple s in S:
    - if h(s) in HT:
      - check forall r where  $h(r) = h(s)$  if  $r = s$ 
        - output
- \* Cost  $\mathcal{O}(B(S) + B(R))$ 
  - Assumed hash table fits into memory
  - Not good when hash table does not fit into memory
- \* **Efficient if:** Result of join fits into memory

◇ **Grace Hash Join**

- + Deals with the case where the hash table does not fit into memory
- \* Idea
  - Partition  $R$  and  $S$  by hashing them using  $h1$ 
    - ▷ Matching tuples of  $R$  and  $S$  are mapped into the same partition
    - ▷ Only one partition of  $R$  and  $S$  into memory and compare them
  - Rehash the hashes using  $h2$  and check for equality
- \* Cost  $\mathcal{O}(3B(R) + 3B(S))$

### 15.5.1 Query Optimizer

- Given a SQL query generate a good execution plan.
  - ◇ **Execution Plan:** Tree of relational algebra operators

- Used query executor to actually execute the relational algebra operators
- Terminology
  - ◊ **Logical Plan:** What the user logically wants
  - ◊ **Physical Plan:** What the DMBS can understand and run
- Steps
  - ◊ User gives SQL query
  - ◊ Parse SQL to logical plan
  - ◊ Convert to physical plan
  - ◊ Run each operator using the operator execution
- **Execution Model:** How different operators are put together
  - ◊ Different ways to put operators together
  - ◊ **Iterator Model**
    - \* Each operator is an iterator
      - **Input:** Set of streams of tuples
      - **Output:** Stream of tuples
      - Calling `next()` on a iterator returns its result
    - \* **Query Plan:** Is tree of iterators
      - Results returned by calling `root.next()` over and over
    - \* **Volcano Model:** Data flow from bottom to top
  - + Generic interface for all operators
  - + Easy to implement iterators
  - + No overheads in terms of main memory
  - + Supports pipelining
  - + Supports parallelism and distribution
    - High overhead of method calls
    - Poor instruction cache locality
  - ◊ **Materialization Model**
    - \* Each operator processes its inputs all at once and then emits its output all at once
      - **Input:** Full relation
      - **Output:** Full relation
    - \* Good when the intermediate result is not too much larger than the final result
      - + Good for OLTP
      - Bad for OLAP
  - ◊ **Vectorization Model:**
    - \* Similar to iterator model
    - \* Each operator returns a batch of tuples
      - Instead of a single tuple
  - + Good for OLAP
  - + Allows for vectored instructions to process batches of tuples
- **Cost Model:** How to estimate the cost of each physical plan given our execution model
  - ◊ Cannot be calculated but only estimate
  - ◊ Estimated based on lower level (operator execution)
  - ◊ Key variable to estimate is selectivity
    - \* Hard to estimate selectivity
  - ◊ **Cardinality Estimation**
    - \* How many tuples does a query involve?
    - \* **Histogram**
      - Create histograms from the data

- Lookup histograms to get an estimate of the number of tuples
  - Hard to combine multiple histograms into one
    - ▷ Due to missing correlation
    - ▷ Assume they are independent
  - For continuous values use bins
  - We can also create multidimensional histograms
- \* Can benefit from machine learning
- **Space Space:** What are the logically equivalent sets of physical plans?
  - ◇ Given an input logical plane, there are different ways that one construct a physical plan
    - \* I.e. different orders of operators
  - ◇ **Query rewriting Rules:** Set of transformations
  - ◇ **Input:** Relational algebra expression  $E$
  - ◇ **Output:** Relational algebra expression  $E'$
  - ◇ **Property:**  $E$  is equivalent to  $E'$ 
    - \*  $\forall I \in \underbrace{\mathbf{I}}_{\text{Set of possible DB instances}}, E(I) = E'(I)$
  - ◇ Many different rules are possible
    - 1) Conjunctive selection operations can be deconstructed into a sequence of individual selections
      - $\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$
    - 2) Selection operations are commutative
      - $\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$
    - 3) Only the last in a sequence of projections operations is needed
      - $\Pi_{t_1}(\Pi_{t_2}(E)) = \Pi_{t_1}(E)$
    - 4) Selections can be combined with Cartesian products and theta joins
      - $\sigma_{\theta}(E_1 \times E_2) = E_1 \bowtie_{\theta} E_2$
      - $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$
    - 5) Theta-join and natural join operations are commutative
      - $E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$
    - 6) Natural join operations are associative
      - $(E_1 \bowtie E_2) \bowtie E_3 = E_1 \bowtie (E_2 \bowtie E_3)$
    - 6b) Theta join operators are associative
      - $(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$  when  $\theta_2$  involves only attributes from  $E_2$  and  $E_3$
    - 7) Pushdown Selection
      - $\sigma_{\theta}(E_1 \bowtie E_2) = \sigma_{\theta}(E_1) \bowtie (E_2)$  if  $\theta$  only involves attributes in  $E_1$
    - 8) The projections operation distributes over the theta join operation
      - $\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) = (\Pi_{L_1}(E_1)) \bowtie_{\theta} (\Pi_{L_2}(E_2))$  if  $\theta$  only involves attributes from  $L_1 \cup L_2$
    - 9) The set operations union and intersection are commutitive
      - $E_1 \cup E_2 = E_2 \cup E_1$
      - $E_1 \cap E_2 = E_2 \cap E_1$
    - 10) Set union and intersection are associative
      - $(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$
      - $(E_1 \cap E_2) \cap E_3 = E_1 \cap (E_2 \cap E_3)$
    - 11) Selection operation distrubites over  $\cup, \cap$  and  $\setminus$ 
      - $\sigma_{\theta}(E_1 \setminus E_2) = \sigma_{\theta}(E_1) \setminus \sigma_{\theta}(E_2) = \sigma_{\theta}(E_1) \setminus E_2$
      - $\sigma_{\theta}(E_1 \cup E_2) = \sigma_{\theta}(E_1) \cup \sigma_{\theta}(E_2) = \sigma_{\theta}(E_1) \cap E_2$

- $\sigma_\theta(E_1 \cap E_2) = \sigma_\theta(E_1) \cap \sigma_\theta(E_2) \neq \sigma_\theta(E_1) \cup E_2$
- 12) Projection operation distributes over union
  - $\Pi_L(E_1 \cup E_2) = (\Pi_L(E_1)) \cup (\Pi_L(E_2))$
- **Search Algorithm:** How can we search the best physical plan, given cost model?
  - ◊ Searching for the optimal query is hard
  - ◊ **Compromise 1:** Constraint search space
    - \* Only consider left-deep join trees
    - \* Allows fully pipelined plans
      - Intermediate values do not need to be written to temporary files
      - Not all trees are fully pipelined (e.g. Sort Merge Join)
    - \* Process:
      - Enumerate join orders
        - ▷ Only different left deep trees
      - Enumerate plans for each operator
      - Enumerate access method for each operator
    - \* Complexity  $\mathcal{O}(n!)$ ,  $n$  is # relations
    - \* Can be done with DP
  - ◊ **Compromise 2:** Heuristic-based Optimisation
    - \* Optimize query-tree by applying a set of rules that typically improve execution performance
      - Early selection
      - Early projection
      - Restrictive selections and joins before other similar operators
    - \* Search algorithm that enables pipelining
- EXPLAIN gives actual physical model

## 15.6 Tuning

- DBMS have many internal parameters
- Improve performance

## 16 Transaction

- **Motivation**
  - ◊ Assume:
    - \* DB is a collection of objects
      - I.e. one tuple is one object
    - \* Objects are fixed
      - We cannot create new ones or delete old ones
    - \* System has only a single CPU
      - CPU can only run one instruction at the time
  - ◊ If not dealt with correctly, simultaneous transactions may get mixed and we get wrong results
  - ◊ **Concurrent DB Access**
    - \* **Schedule:** One way of mixing instructions
      - Different schedules may result in different results
    - \* Result of one query may be overwritten partly or completely
    - \* **Attribute-level Inconsistency:** Concurrent change of a single attribute of the same tuple
    - \* **Tuple-level Inconsistency:** Concurrent change of different attributes of the same tuple
    - \* **Table-level Inconsistency:** Concurrent change of full relation
    - \* **Multi-statement Inconsistency:** Interleaving of concurrent queries
    - \* When multiple groups of SQL statements are running at the same time, we want the effect as if they are executed sequentially
  - ◊ **System Failure**
    - \* Many thing which can break in a real system
    - \* We want that all or none changes apply, but not partial application
- **Transaction:** Collection of instructions which should not mix with other transactions
  - ◊ Concurrent transactions appear to run in isolation
  - ◊ On a crash, transaction changes appear entirely or not at all
  - ◊ BEGIN; ..... COMMIT: Encapsulates a transaction
    - \* Transaction has finished, database confirms to client whan all changes of the transaction have been made persistent
    - \* Transaction may also fail. Database rollback all changes done by the transaction
      - Written as BEGIN; ..... ABORT;
      - Can be initiated by the user or DBMS
  - ◊ **Autocommit:** If true, turns each SQL statement into own transition
    - \* SQL option
    - \* Activated by default, can be deactivated

### 16.1 ACID:

- Desired properties of transaction
- **Atomicity**
  - ◊ Transaction is executed in its entirety or not at all
- **Consistency**
  - ◊ A committed transaction goes from one consistent state to another consistent state
    - \* Before and after a transaction all integrity constraints must hold
  - ◊ Within a transaction constraints may be violated
  - ◊ Transaction leads from consistent state to consistent state



- ◇ Granularity depends on the integrity constraints
  - \* I.e. some constraints are checked for each tuple, some for each statement and some for a transaction
  - \* Can be controlled and influenced to some degree
- **Isolation**
  - ◇ **Ideally:** Transaction executes as if it were alone in the system
    - \* I.e. enforce serializability
    - \* Much too hard to enforce
  - ◇ Implies that integrity constraints always hold if each transaction is correct
  - ◇ DMDB picks one execution order at random (if there are multiple)
    - \* If not desired, application must enforce this
- **Durability**
  - ◇ If system crashes after a transaction, the changes of the transaction must still remain in the DB
    - \* Or somehow recoverable

## 16.2 Isolation

- One of the key properties
- **Anomaly:** Misbehaviour of the DB
  - ◇ **Dirty Read:** Read a value which was updated by another transaction which has not yet committed
    - \* May contain values which were/are never in the DB
      - When the other transaction aborts
  - ◇ **Non-repeatable Reads:** Reading the same tuple twice gives give us different values both times
    - \* It was updated by another transaction which committed (difference to dirty read)
  - ◇ **Phantoms:** During a transaction, another transaction added or removed tuples
    - \* Similar to non-repeatable reads
- **Isolation Level:** Defines for each transaction what anomalies we allow to happen

	Dirty Reads	Non-Repeatable Reads	Phantoms	overhead ↓ ↑ concurrency
Read Uncommitted	✓	✓	✓	
◇ Read Committed	×	✓	✓	
Repeatable Read	×	×	✓	
Serializable	×	×	×	

## 16.3 More on Serializable

- **Serializable:** Schedule that leads to the same answer as some serial schedule
  - ◇ Only depends on final result and not I/O pattern along
  - ◇ Not all sequential orders necessarily lead to the same result
    - Hard or impossible to enforce
- **Conflicts**
  - ◇ **Definition**
    - \* **Same Transaction:**
      - Two operations are always conflicting
      - $\implies$  Reordering within transaction is not allowed
    - \* **Different Transaction  $O_1$  in  $T_1$  and  $O_2$  in  $T_2$ :**
      - $O_1$  and  $O_2$  are conflicting if one of them is a write to the same location
  - ◇ **Types**

- \* **Read-Write:**
  - Leads to unrepeatable reads
- \* **Write-Read:**
  - Leads to dirty read
- \* **Write-Write:**
  - Leads to overwriting of uncommitted data
- **Conflict Equivalent** are two schedules iff:
  - One can be transformed into the other by swapping non-conflicting operations
- **Conflict Serializable:** Schedule if it is conflict equivalent to some serial schedule
  - I.e. schedule which can be translated into a serial schedule with a sequence of non-conflicting swaps of adjacent actions
  - Stronger than serializable
    - \* Conflict serializable  $\subseteq$  serializable
  - Only depends on the read/write pattern
    - \* And not what we are writing
  - Easier than serializability for DB to handle as it does not require the DB to understand what each operator is doing
  - Enforced by most DBMSs
  - **Decide**
    - \* Each transaction is a node
    - \*  $\exists$  edge  $T_i$  to  $T_j$  if:
      - Operator  $o_i$  in  $T_i$  is in conflict with operator  $o_j$  in  $T_j$
      - $o_j$  appears earlier than  $o_i$  in same transaction
    - \* Schedule is conflict serializable iff its dependency graph is acyclic
- Serializability and conflict serializability only concern committed transactions (and not aborted)
  - Operations can be in conflict with ABORT
    - \* E.g. `READ(X)` / `WRITE(X)` before or after `ABORT` of other transaction may lead to different results (only if the other transaction did `WRITE(X)`)
  - But they are somehow still considered as serializable

## 16.4 Enforce Isolation

- **Goal:** Only allow schedules that are conflict serializable
- Two main approaches
  - **Pessimistic:** Assume that conflicts happen all the time
    - \* Use locks
  - **Optimistic:** Assume that most transaction do not conflicts
    - \* Use snapshot isolation
- Need to evaluate which approach is better for our application

### 16.4.1 Locking

- **Assume:** Do not know what transactions are going to do in the future
- **Idea:** Before the system access the data object  $X$  it locks  $X$ 
  - Prevents access of  $X$  by other transaction
- Lock is released only when it is safe to
  - I.e. the execution is guaranteed to be conflict serializable
- Allows to enforce serializable schedule
- Types

- ◇ **Shared Lock (S Lock):**
    - \* For reading
  - ◇ **Exclusive Lock (X Lock):**
    - \* For writing
- Does not necessarily enforce conflict serializability
- **Two-Phase Locking (2PL)**
  - ◇ Consists of two phases
    - \* **Phase 1: Growing**
      - Acquire required locks
      - Cannot release any locks
    - \* **Phase 2: Shrinking**
      - Release locks
      - Cannot acquire new locks
- + Guarantees conflict serializability
  - **Cascading Abort:** Abort of one transaction leads to abort of another transaction
    - \* Happens when we read from a transaction which gets aborted
    - \* Really bad if we have already committed
      - Commit must be undone
      - Conflicts with durability property
- **Strict Two-Phase Locking (Strict 2PL)**
  - ◇ **Phase 1:** is similar to 2PL
  - ◇ **Phase 2:**
    - \* All locks are kept until end of transaction
      - I.e. COMMIT or ABBORT released all locks
  - Deadlocks possible
    - \* **Detection**
      - Each transaction is a node
      - $\exists$  edge  $T_i$  to  $T_j$  if  $T_i$  is waiting for a lock currently hold by  $T_j$
    - \* Deadlock if we have a cyclic wait-for graph
    - \* Non-trivial to decide which transaction to kill
    - \* Prevented by locking in some global order
- **Granularity**
  - ◇ Problems
    - \* We need to lock every single tuple at its own
    - \* We need to hold locks for whole transaction to prevent phantoms
  - ◇ DB is hierarchical structure and hence needs to support hierarchical locking
  - ◇ New locks
    - \* **Intention Share (IS):**
      - Some lower nodes are in S
    - \* **Intention Exclusive (IX):**
      - Some lower node are in X
    - \* **Share and Intention Exclusive (SIX):**
      - Root is locked in S
      - Some lower node are in X
  - ◇ Old locks
    - \* **S:** All lower nodes are in shard
    - \* **X:** All lower nodes are in exclusive
  - ◇ Full overview

	Mode	Current Lock					
		NL	IS	IX	S	SIX	X
Request	NL	✓	✓	✓	✓	✓	✓
	IS	✓	✓	✓	✓	✓	×
	IX	✓	✓	✓	×	×	×
	S	✓	✓	×	✓	×	×
	SIX	✓	✓	×	×	×	×
	X	✓	×	×	×	×	×

◇ Steps to lock a tuple

- 1) Acquire *IS* on database
- 2) Acquire *IS* on table
- 3) Acquire *S* on tuple

### 16.4.2 Snapshot Isolation

- **Idea:** Assume that transaction are serializable and revert if they are not
- Working
  - ◇ Transaction receives timestamp  $TS(T)$  when it starts
  - ◇ Reads are carried out as of the DB version of  $TS(T)$
  - ◇ Writes are carried out on a separate buffer
  - ◇ When transaction commits, abort  $T_1$  if  $\exists T_2$  such that:
    - \*  $T_1$  and  $T_2$  update the same object and
    - \*  $T_2$  committed after  $TS(T_1)$  but before  $T_1$  commits
  - ◇ Instead of aborting  $T_1$ , we can also let  $T_1$  finish and only then merge  $T_2$
- **Timestamps**
  - ◇ System clock or monotonically increasing for each transaction
- + High concurrency and availability
  - ◇ Only block when transaction commits
- + No cascading abort
- + No deadlock
- Unnecessary rollbacks
- **Write Skew:** Interaction of multiple objects
  - ◇ Checking integrity constrains happens in the snapshots
  - ◇ Two concurrent transaction update different objects
  - ◇ Integrity constrains for each is ok but for the combination not
- Looser version
  - ◇ Idea
    - \* Object themselves have read (last read) and write (last written) timestamps
    - \* If transaction accesses object from the future (object has higher timestamp than transaction timestamp), transaction is aborted
    - \* Object timestamps are updated on read/writes
- + No unnecessary rollbacks (I guess)
  - Long transaction may starve
  - Cascading aborts

## 17 Recoverability

### 17.1 Definition

- Important for durability property
- Ensures that the state in the DB is correct
- Need to recover from
  - ◊ Abort of single transaction
    - \* Undo all changes of the aborted transaction
  - ◊ System crash (loss of main memory but not disk)
    - \* Redo all committed transaction
- **T<sub>1</sub> reads from T<sub>2</sub>:** if T<sub>1</sub> reads a value written by T<sub>2</sub> at a time when T<sub>2</sub> was not aborted
- Different families of schedules
  - ◊ Each family has different recoverability properties
- **Recoverable (RC):**
  - ◊ If T<sub>i</sub> reads from T<sub>j</sub> and commits, then  $c_j < c_i$
  - ◊ If T<sub>i</sub> reads from T<sub>j</sub> and aborts, or if T<sub>i</sub> writes etc. it is also RC
  - ◊ No need to undo a committed transaction
  - ◊ **If not RC:** Loss of data
- **Avoids Cascading Aborts (ACA):**
  - ◊ If T<sub>i</sub> reads X from T<sub>j</sub>, then  $c_j < r_i[X]$
  - ◊ If T<sub>i</sub> writes X from T<sub>j</sub>, it is also ACA
  - ◊ Aborting a transaction does not cause aborting others
  - ◊ **If not ACA:** Thrashing behaviour when transactions abort each other
- **Strict (ST):**
  - ◊ If T<sub>i</sub> reads from or writes a value written by T<sub>j</sub>, then
    - \* **If T<sub>j</sub> commits:**  $(c_j < r_i[X] \wedge c_j < w_i[X])$
    - \* **If T<sub>j</sub> aborts:**  $(a_j < r_i[X] \wedge a_j < w_i[X])$
  - ◊ Extends ACA to write
  - ◊ Undoing a transaction does not undo the changes of other transactions
  - ◊ **If not ST:** Recovery is very complex or impossible
  - ◊ Enforced by Strict 2PL
- All Schedules  $\subset$  RC  $\subset$  ACA  $\subset$  ST  $\subset$  Serial
- **Goal:** All allowed schedules lie in the intersection of ST and conflict serializable

### 17.2 Write-Ahead Log

- **Assume:**
  - ◊ Disk is save
  - ◊ **Write(A, v)** only changes object in memory but not disk
  - ◊ **OUTPUT(A)** writes changes from memory to disk
- **Idea:** Log changes and restore from log if required
- **Log:** File which only can be appended to
  - ◊ Stored in memory and periodically flashed to disk
    - \* When to flash?
  - ◊ Operation
    - \* Append Record
      - START T
      - COMMIT T
      - ABORT T

- Update  $\langle T, X, v \rangle$
  - \* Flush to disk
    - FLUSH
  - \* Log message do not have to mean that the action was actually done on the DB
- Two main strategies
- **Undo Logging**
  - ◊ If  $T$  modified DB element  $X$  log  $\langle T, X, \text{old value} \rangle$  to disk before change  $X$  is written to disk
    - \* I.e. call FLUSH before calling OUTPUT
  - ◊ If a transaction commits, COMMIT record must be logged to disk only after all other changes are written to disk
    - \* I.e. log COMMIT and call FLUSH only after calling OUTPUT
  - ◊ Recovery
    - \* **Committed Transaction:** Ones which have COMMIT in the log
      - COMMIT in log guarantees that changes are flushed to disk
        - ▷ Nothing to do
    - \* **Uncommitted Transaction:** Ones which do not have COMMIT in the log
      - We cannot be sure if changes were committed or not
        - ▷ Undo everything
      - Steps
        - ▷ Find all transaction  $i$  with Start  $T_i$  but not COMMIT  $T_i$
        - ▷ If there is only a single transaction:
          - Scan from the end and undo updates
        - ▷ If there are multiple uncommitted transactions
          - Scan from the end (skipping logs from committed transaction) and undo updates
        - ▷ Write ABORT  $T_i$  at the end of log
        - ▷ Flush log
  - Lots of I/O
  - Log is almost the size of the transaction
  - **Redo Logging**
    - ◊ If  $T$  modifies DB element  $X$  log  $\langle T, X, \text{new value} \rangle$  to disk
    - ◊ Log COMMIT and call FLUSH, before calling OUTPUT
    - ◊ Recovery
      - \* Scan log from the beginning
      - \* If COMMIT  $T_i$  is not in the log
        - No changes of  $T_i$  appears on disk
        - Write ABORT  $T_i$
        - Ignore changes if  $T_i$  during scanning
      - \* If COMMIT  $T_i$  is in the log
        - Does not mean all its changes are already on disk
        - Redo all changes of  $T_i$
      - \* Flush log
  - + Less I/O than undo logging (I guess)
    - The log we need to keep can be very, very long
      - \* After commit we still do not know if the changes are reflected on the DB
    - Problematic when two transaction update different objects which are stored on the same page
  - **Undo/Redo Logging**

- ◇ Combine both to get pro from both
- ◇ Before modifying any DB element  $X$  on disk, write  $\langle T, X, \text{old value}, \text{new value} \rangle$
- ◇ Flush log before actual changes are made on disk
- ◇ Recovery
  - \* If **COMMIT**  $T$  is not in the log
    - $T$  is incomplete
    - Undo changes
  - \* If **COMMIT**  $T$  is in the log
    - $T$  is complete
    - Redo changes

## 18 Distribution

### • Distributed Commit

- ◇ **Problem:** Multiple DBs and a single coordinator which manages them
  - \* Atomicity of a single nodes does not imply atomicity in a distributed setting
- ◇ Two main methods
- ◇ **Two Phase Commit**
  - \* Consists of two phases
  - \* **Voting Phase:** Coordinator inquires if all nodes are ready and willing to commit
    - Initiated by coordinator sending **Prepare**
    - If node says *OK* it cannot change its mind anymore
  - \* **Decision Phase:** Coordinator asks all nodes to commit
    - Initiated by coordinator sending **Commit**
    - Coordinator sends **Abort** if not all workers are ready
    - Only if all said *OK* in voting phase
    - If any worker or coordinator dies we have to rollback
  - \* If a worker/coordinator is temporarily dead we may continue but let the other know that we were dead
- ◇ **Linear Two Phase Commit**
  - \* Coordinator only communicates with one workers, which communicates with another worker etc...
- ◇ **Two Phase Commit vs Linear Two Phase Commit**

	Two Phase Commit	Linear Two Phase Commit
* Given 1 Coordinator and $N$ Workers		
* Total Messages	$3N$	$2N$
* Latency $t$	$3t$	$2Nt$

### • Distributed Query Processing

- ◇ Execute query on multiple machines
  - \* Desirable for
    - Data too large to fit into one machine
    - Computationally intensive query
- ◇ Different ways of construction
  - \* Shared Memory
  - \* Shard Disk
  - \* Nothing Shared
    - We consider this one
    - Master received query and distributes to several workers
- ◇ **Goal:** Hide complexity from users
- ◇ **Idea:** Partition DB to each worker and each only deals with own partition
- ◇ Examples
  - \* **Table Partitioning**
    - ```
SELECT * FROM R(a,b,c), S(a,d)
WHERE R.b = 1
AND S.d = 2
AND R.a = S.a
```
    - Worker 1
      - ▷ Run
 

```
SELECT * FROM R
WHERE R.b = 1
```



- ▷ Generate table  $R'(a, b, c)$
  - Worker 2
    - ▷ Run
 

```
SELECT * FROM S
WHERE S.d = 2
```
    - ▷ Generate table  $S'(a, d)$
    - Combine  $R'$  and  $S'$ 
      - ▷
 

```
SELECT * FROM R', S'
WHERE R'.a = S'.a
```
- + Worker 1 and 2 can work concurrency
- + If  $R'$  and  $S'$  are small the communication is small
  - Problematic if one table is too large for a worker
  - Parallelizability of queries (if available) is not exploited
- \* **Horizontal Partitioning**
  - ```
SELECT * FROM R(a, b, c), S(a, d)
AND R.a = S.a
```
  - Worker 1
    - ▷ Run
 

```
SELECT * FROM R1, S1
WHERE R1.a = S1.a
```
    - ▷ Generate table  $T1(a, b, c, d)$
    - Worker 2
      - ▷ Run
 

```
SELECT * FROM R2, S2
WHERE S2.a = R2.a
```
      - ▷ Generate table  $T2(a, b, c, d)$
      - Combine  $T1$  and  $T2$ 
        - ▷
 

```
SELECT * FROM T1
UNION
SELECT * FROM T2
```
  - + When the result is small this can be fast TODO: Not sure why this even works
  - \* **Distributed QO 1**
    - ```
SELECT * FROM R(a, b, c), S(a, d)
WHERE R.a = S.a
```
    - Replicate one table on both nodes and split the other in two
    - Worker 1:  $T_1 = R_1 \bowtie S$
    - Worker 2:  $T_2 = R_2 \bowtie S$
    - Combine:  $T_1 \cup T_2$
  - \* **Distributed QO 2**
    - ```
SELECT * FROM R(a, b, c), S(a, d)
```

WHERE R.a = S.a

- Tables are portioned on the join attribute and each node performs the join locally
- Worker 1:  $T_1 = R_1 \bowtie S_1$
- Worker 2:  $T_2 = R_2 \bowtie S_2$
- Combine:  $T_1 \cup T_2$

\* **Distributed QO 3**

- ```
SELECT * FROM R(a, b, c), S(a, d)
WHERE R.a = S.a
```

- Tables are portioned on different keys
- Worker 1:  $T_1 = R_1 \bowtie (S_1 \cup S_2)$
- Worker 2:  $T_2 = R_2 \bowtie (S_2 \cup S_2)$
- Combine:  $T_1 \cup T_2$

◇ **Replication**

- \* Replicate data among several machines
- \* **Group Mirroring**
  - Replicate at machine level
  - Each machine has data blocks which are stored on two other machines
    - If the two wrong machines die we loose data
    - Death of a single machine leads to  $2x$  slowdown
- \* **Spread Mirroring**
  - Each machine contains data available on all other machines
    - If two die we certainly loose (little) data
    - Single failure leads to  $1/N$  more load on other machines

● **Distributed Key-Value Store**

- ◇ Relational DB is expensive and does not scale well
- ◇ Data Model
  - \* Key + Value
  - \* Indexed on key
- ◇ Distributed Deployment
  - \* Horizontal partitioning
  - \* Replication of partitions
- ◇ Build
  - \* Build as a simple hash table
  - \* If distributed, we use consistent hashing
- + Very fast lookups
- + Easy to scale
  - \* Add more copies as we add more machines
- Only support point queries
- Some operations are very expensive
- Hard to keep data consistent among different copies
- Complexity is pushed to user/application